

IMPORTING PACKAGES

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
#from sklearn.feature_selection import SelectKBest,f_classif,mutual_info_classif
from sklearn.model_selection import cross_validate,train_test_split
from sklearn.metrics import classification_report,confusion_matrix
from sklearn.naive_bayes import GaussianNB
#from scikitplot.metrics import plot_confusion_matrix
#from sklearn.metrics import plot_confusion_matrix
from sklearn.linear_model import LogisticRegression
import seaborn as sns
```

LOADING DATASET

```
In [4]: data=pd.read_csv("C:/Users/DELL/Downloads/creditcard.csv")
data
```

Out[4]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012468
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777

284807 rows x 31 columns

```
In [5]: data.head()
```

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267

5 rows x 31 columns

DATA EXPLORATION

In [6]: `data.shape`

Out[6]: (284807, 31)

In [7]: `data.describe()`

Out[7]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.919560e-15	5.688174e-16	-8.769071e-15	2.782312e-15	-1.552563e-15	2.010663e-15	-1.694249e-15	-1.927028e-16	-3.137024e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+02
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+02

8 rows x 31 columns

In [8]: `#printing the columns`

```
print('Columns Name:',list(data))
print('Total Number of Columns: ',len(list(data)))
```

Columns Name: ['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount', 'Class']
Total Number of Columns: 31

In [9]: `#counting the number of genuine and fraudulent Transactions`

```
n_genuine=data[data['Class']==0]
n_fraud=data[data['Class']==1]
```

`#displaying the number of genuine and fraudulent Transactions`

```
print('Genuine Transactions: ',len(n_genuine))
print('Fraudulent Transactions: ',len(n_fraud))
```

Genuine Transactions: 284315

Fraudulent Transactions: 492


In [10]: `outlierfraction=float(len(n_fraud))/float(len(n_genuine))`

```
print(outlierfraction)
```

0.0017304750013189597

In [11]: `#plotting the pie plot of genuine and fraudulent transactions`

```
plt.pie([len(n_genuine),len(n_fraud)],labels=['Genuine','fraud'],radius=1)
plt.show()
```



```
In [12]: n_fraud.Amount.describe()
```

```
Out[12]: count    492.000000
         mean     122.211321
         std      256.683288
         min       0.000000
         25%       1.000000
         50%       9.250000
         75%      105.890000
         max      2125.870000
         Name: Amount, dtype: float64
```

```
In [13]: n_genuine.Amount.describe()
```

```
Out[13]: count    284315.000000
         mean      88.291022
         std      250.105092
         min       0.000000
         25%       5.650000
         50%      22.000000
         75%      77.050000
         max     25691.160000
         Name: Amount, dtype: float64
```

SEPERATING DATA INTO X AND Y

```
In [14]: #here we put training data in x variable and target value in y variable
         x,y=data.iloc[:, :-1],data.iloc[:, -1]
```

```
In [15]: x.head()
```

```
Out[15]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	0.251412	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.069083	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.524980	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.208038	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	0.408542	-0.009431	0.798278	-0.137458

5 rows x 30 columns

```
In [16]: y.head()
```

```
Out[16]: 0    0
         1    0
         2    0
         3    0
         4    0
         Name: Class, dtype: int64
```

SPLITTING OF DATA INTO TRAINING AND TESTING DATA

```
In [17]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=42)
```

Building a Random Forest Classifier

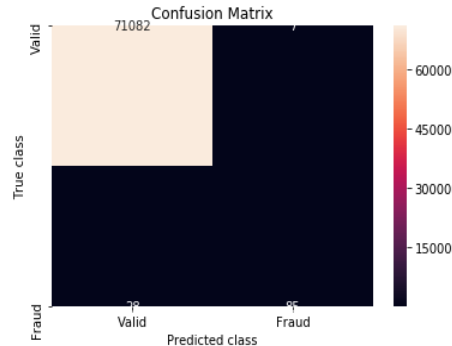
```
In [18]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
y_pred=rfc.predict(x_test)
```

C:\Users\DELL\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:245: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

EVALUATING THE MODEL

```
In [19]: labels=['Valid','Fraud']
conf_matrix=confusion_matrix(y_test,y_pred)

sns.heatmap(conf_matrix,xticklabels=labels,yticklabels=labels,annot=True,fmt='d')
plt.title("Confusion Matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()
```



```
In [20]: from sklearn.metrics import classification_report,accuracy_score,precision_score,recall_score,f1_score
print("The model used is random Forest classifier")
acc=accuracy_score(y_test,y_pred)
print("The accuracy Score is {}".format(acc))
prec=precision_score(y_test,y_pred)
print("The precision Score is {}".format(prec))
rec=recall_score(y_test,y_pred)
print("The recall Score is {}".format(rec))
f1=f1_score(y_test,y_pred)
print("The f1 Score is {}".format(f1))
```

The model used is random Forest classifier
The accuracy Score is 0.9995084407741356
The precision Score is 0.9239130434782609
The recall Score is 0.7522123893805309
The f1 Score is 0.8292682926829267

LOGISTIC REGRESSION


```
In [21]: lr=LogisticRegression()  
lr.fit(x_train,y_train)
```

```
C:\Users\DELL\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.  
FutureWarning)
```

```
Out[21]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
intercept_scaling=1, l1_ratio=None, max_iter=100,  
multi_class='warn', n_jobs=None, penalty='l2',  
random_state=None, solver='warn', tol=0.0001, verbose=0,  
warm_start=False)
```

```
In [22]: pred=lr.predict(x_test)
```

```
In [23]: labels=['Valid','Fraud']  
conf_matrix=confusion_matrix(y_test,pred)  
  
sns.heatmap(conf_matrix,xticklabels=labels,yticklabels=labels,annot=True,fmt='d')  
plt.title("Confusion Matrix")  
plt.ylabel("True class")  
plt.xlabel("Predicted class")  
plt.show()
```



```
In [24]: from sklearn.metrics import classification_report,accuracy_score,precision_score,recall_score,f1_score  
print("The model used is logistic regression classifier")  
acc=accuracy_score(y_test,pred)  
print("The accuracy Score is {}".format(acc))  
prec=precision_score(y_test,pred)  
print("The precision Score is {}".format(prec))  
rec=recall_score(y_test,pred)  
print("The recall Score is {}".format(rec))  
f1=f1_score(y_test,pred)  
print("The f1 Score is {}".format(f1))
```

```
The model used is logistic regression classifier  
The accuracy Score is 0.9989607033510295  
The precision Score is 0.7294117647058823  
The recall Score is 0.5486725663716814  
The f1 Score is 0.6262626262626262
```

LINEAR SVM

```
In [25]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
```

```
In [26]: svm_clf=LinearSVC()
```

```
In [27]: svm_clf.fit(x_train,y_train)
```


C:\Users\DELL\Anaconda3\lib\site-packages\sklearn\svm\base.py:929: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)

```
Out[27]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0)
```

```
In [28]: y_predi=svm_clf.predict(x_test)
```

```
In [29]: labels=['Valid','Fraud']
conf_matrix=confusion_matrix(y_test,y_predi)

sns.heatmap(conf_matrix,xticklabels=labels,yticklabels=labels,annot=True,fmt='d')
plt.title("Confusion Matrix")
plt.ylabel("True class")
plt.xlabel("Predicted class")
plt.show()
```



```
In [30]: from sklearn.metrics import classification_report,accuracy_score,precision_score,recall_score,f1_score
print("The model used is LinearSVC")
acc=accuracy_score(y_test,y_predi)
print("The accuracy Score is {}".format(acc))
prec=precision_score(y_test,y_predi)
print("The precision Score is {}".format(prec))
rec=recall_score(y_test,y_predi)
print("The recall Score is {}".format(rec))
f1=f1_score(y_test,y_predi)
print("The f1 Score is {}".format(f1))
```

```
The model used is LinearSVC
The accuracy Score is 0.9985112777730962
The precision Score is 0.64
The recall Score is 0.1415929203539823
The f1 Score is 0.2318840579710145
```

```
In [31]: Algorithms=['RandomForest','LogisticRegression','LinearSVM']
Accuracy=[0.995,0.9987,0.9984]
Precision=[0.946,0.611,0.6190]
recall=[0.786,0.584,0.1150]
final=pd.DataFrame({'Accuracy':Accuracy,'Algorithm':Algorithms,'Precision':Precision,'recall':recall})
final
```

Out[31]:

	Accuracy	Algorithm	Precision	recall
0	0.9950	RandomForest	0.946	0.786
1	0.9987	LogisticRegression	0.611	0.584
2	0.9984	LinearSVM	0.619	0.115

```
In [32]: sns.lineplot(x='Algorithm',y='Accuracy',data=final,palette='hot',label='Accuracy')
sns.lineplot(x='Algorithm',y='Precision',data=final,palette='hot',label='Precision')
sns.lineplot(x='Algorithm',y='recall',data=final,palette='hot',label='recall')
```

Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x17f2924c8c8>

